



[About Keras](#)

[Getting started](#)

[Developer guides](#)

Keras 3 API documentation

Models API

Layers API

Callbacks API

Ops API

Optimizers

Metrics

Losses

Data loading

Built-in small datasets

Keras Applications

Xception

EfficientNet B0 to B7

EfficientNetV2 B0 to B3 and S, M, L

ConvNeXt Tiny, Small, Base, Large, XLarge

VGG16 and VGG19

ResNet and ResNetV2

MobileNet, MobileNetV2, and
MobileNetV3

DenseNet

NasNetLarge and NasNetMobile

InceptionV3

InceptionResNetV2

Mixed precision

Multi-device distribution

RNG API

Utilities

KerasTuner

KerasCV

KerasNLP

[Keras 2 API documentation](#)

[Code examples](#)

[KerasTuner: Hyperparameter Tuning](#)

[KerasCV: Computer Vision Workflows](#)

MobileNet, MobileNetV2, and MobileNetV3

MobileNet function

[\[source\]](#)

```
keras.applications.MobileNet(  
    input_shape=None,  
    alpha=1.0,  
    depth_multiplier=1,  
    dropout=0.001,  
    include_top=True,  
    weights="imagenet",  
    input_tensor=None,  
    pooling=None,  
    classes=1000,  
    classifier_activation="softmax",  
    name=None,  
)
```

Instantiates the MobileNet architecture.

Reference

- [MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications](#)

This function returns a Keras image classification model, optionally loaded with weights pre-trained on ImageNet.

For image classification use cases, see [this page for detailed examples](#).

For transfer learning use cases, make sure to read the [guide to transfer learning & fine-tuning](#).

Note: each Keras Application expects a specific kind of input preprocessing. For MobileNet, call `keras.applications.mobilenet.preprocess_input` on your inputs before passing them to the model. `mobilenet.preprocess_input` will scale input pixels between -1 and 1.

Arguments

- **input_shape**: Optional shape tuple, only to be specified if `include_top` is `False` (otherwise the input shape has to be `(224, 224, 3)` (with `"channels_last"` data format) or `(3, 224, 224)` (with `"channels_first"` data format). It should have exactly 3 inputs channels, and width and height should be no smaller than 32. E.g. `(200, 200, 3)` would be one valid value. Defaults to `None`. `input_shape` will be ignored if the `input_tensor` is provided.
- **alpha**: Controls the width of the network. This is known as the width multiplier in the MobileNet paper.
 - If `alpha < 1.0`, proportionally decreases the number of filters in each layer.
 - If `alpha > 1.0`, proportionally increases the number of filters in each layer.
 - If `alpha == 1`, default number of filters from the paper are used at each layer. Defaults to `1.0`.
- **depth_multiplier**: Depth multiplier for depthwise convolution. This is called the resolution multiplier in the MobileNet paper. Defaults to `1.0`.
- **dropout**: Dropout rate. Defaults to `0.001`.
- **include_top**: Boolean, whether to include the fully-connected layer at the top of the network. Defaults to `True`.
- **weights**: One of `None` (random initialization), `"imagenet"` (pre-training on ImageNet), or the path to the weights file to be loaded. Defaults to `"imagenet"`.
- **input_tensor**: Optional Keras tensor (i.e. output of `layers.Input()`) to use as image input for the model. `input_tensor` is useful for sharing inputs between multiple different networks. Defaults to `None`.
- **pooling**: Optional pooling mode for feature extraction when `include_top` is `False`.
 - `None` (default) means that the output of the model will be the 4D tensor output of the last convolutional block.

- **avg** means that global average pooling will be applied to the output of the last convolutional block, and thus the output of the model will be a 2D tensor.
 - **max** means that global max pooling will be applied.
- **classes**: Optional number of classes to classify images into, only to be specified if **include_top** is **True**, and if no **weights** argument is specified. Defaults to **1000**.
- **classifier_activation**: A **str** or callable. The activation function to use on the "top" layer. Ignored unless **include_top=True**. Set **classifier_activation=None** to return the logits of the "top" layer. When loading pretrained weights, **classifier_activation** can only be **None** or **"softmax"**.
- **name**: String, the name of the model.

Returns

A model instance.

MobileNetV2 function

[\[source\]](#)

```
keras.applications.MobileNetV2(  
    input_shape=None,  
    alpha=1.0,  
    include_top=True,  
    weights="imagenet",  
    input_tensor=None,  
    pooling=None,  
    classes=1000,  
    classifier_activation="softmax",  
    name=None,  
)
```

Instantiates the MobileNetV2 architecture.

MobileNetV2 is very similar to the original MobileNet, except that it uses inverted residual blocks with bottlenecking features. It has a drastically lower parameter count than the original MobileNet. MobileNets support any input size greater than 32 x 32, with larger image sizes offering better performance.

Reference

- [MobileNetV2: Inverted Residuals and Linear Bottlenecks](#) (CVPR 2018)

This function returns a Keras image classification model, optionally loaded with weights pre-trained on ImageNet.

For image classification use cases, see [this page for detailed examples](#).

For transfer learning use cases, make sure to read the [guide to transfer learning & fine-tuning](#).

Note: each Keras Application expects a specific kind of input preprocessing. For MobileNetV2, call **keras.applications.mobilenet_v2.preprocess_input** on your inputs before passing them to the model. **mobilenet_v2.preprocess_input** will scale input pixels between -1 and 1.

Arguments

- **input_shape**: Optional shape tuple, only to be specified if **include_top** is **False** (otherwise the input shape has to be **(224, 224, 3)** (with **"channels_last"** data format) or **(3, 224, 224)** (with **"channels_first"** data format). It should have exactly 3 inputs channels, and width and height should be no smaller than 32. E.g. **(200, 200, 3)** would be one valid value. Defaults to **None**. **input_shape** will be ignored if the **input_tensor** is provided.
- **alpha**: Controls the width of the network. This is known as the width multiplier in the MobileNet paper.
 - If **alpha < 1.0**, proportionally decreases the number of filters in each layer.
 - If **alpha > 1.0**, proportionally increases the number of filters in each layer.
 - If **alpha == 1**, default number of filters from the paper are used at each layer. Defaults to **1.0**.
- **include_top**: Boolean, whether to include the fully-connected layer at the top of the network. Defaults to **True**.
- **weights**: One of **None** (random initialization), **"imagenet"** (pre-training on ImageNet), or the path to the weights file to be loaded. Defaults to **"imagenet"**.

- **input_tensor**: Optional Keras tensor (i.e. output of `layers.Input()`) to use as image input for the model. `input_tensor` is useful for sharing inputs between multiple different networks. Defaults to `None`.
- **pooling**: Optional pooling mode for feature extraction when `include_top` is `False`.
 - `None` (default) means that the output of the model will be the 4D tensor output of the last convolutional block.
 - `avg` means that global average pooling will be applied to the output of the last convolutional block, and thus the output of the model will be a 2D tensor.
 - `max` means that global max pooling will be applied.
- **classes**: Optional number of classes to classify images into, only to be specified if `include_top` is `True`, and if no `weights` argument is specified. Defaults to `1000`.
- **classifier_activation**: A `str` or callable. The activation function to use on the "top" layer. Ignored unless `include_top=True`. Set `classifier_activation=None` to return the logits of the "top" layer. When loading pretrained weights, `classifier_activation` can only be `None` or `"softmax"`.
- **name**: String, the name of the model.

Returns

A model instance.

MobileNetV3Small function [\[source\]](#)

```
keras.applications.MobileNetV3Small(  
    input_shape=None,  
    alpha=1.0,  
    minimalistic=False,  
    include_top=True,  
    weights="imagenet",  
    input_tensor=None,  
    classes=1000,  
    pooling=None,  
    dropout_rate=0.2,  
    classifier_activation="softmax",  
    include_preprocessing=True,  
    name="MobileNetV3Small",  
)
```

Instantiates the MobileNetV3Small architecture.

Reference

- [Searching for MobileNetV3](#) (ICCV 2019)

The following table describes the performance of MobileNets v3:

MACs stands for Multiply Adds

Classification Checkpoint	MACs(M)	Parameters(M)	Top1 Accuracy	Pixel1 CPU(ms)
mobilenet_v3_large_1.0_224	217	5.4	75.6	51.2
mobilenet_v3_large_0.75_224	155	4.0	73.3	39.8
mobilenet_v3_large_minimalistic_1.0_224	209	3.9	72.3	44.1
mobilenet_v3_small_1.0_224	66	2.9	68.1	15.8
mobilenet_v3_small_0.75_224	44	2.4	65.4	12.8
mobilenet_v3_small_minimalistic_1.0_224	65	2.0	61.9	12.2

For image classification use cases, see [this page for detailed examples](#).

For transfer learning use cases, make sure to read the [guide to transfer learning & fine-tuning](#).

Note: each Keras Application expects a specific kind of input preprocessing. For MobileNetV3, by default input preprocessing is included as a part of the model (as a `Rescaling` layer), and thus `keras.applications.mobilenet_v3.preprocess_input` is actually a pass-through function. In this use case, MobileNetV3 models expect their inputs to be float tensors of pixels with values in the `[0-255]`

range. At the same time, preprocessing as a part of the model (i.e. **Rescaling** layer) can be disabled by setting **include_preprocessing** argument to **False**. With preprocessing disabled MobileNetV3 models expect their inputs to be float tensors of pixels with values in the **[-1, 1]** range.

Arguments

- **input_shape**: Optional shape tuple, to be specified if you would like to use a model with an input image resolution that is not **(224, 224, 3)**. It should have exactly 3 inputs channels. You can also omit this option if you would like to infer input_shape from an input_tensor. If you choose to include both input_tensor and input_shape then input_shape will be used if they match, if the shapes do not match then we will throw an error. E.g. **(160, 160, 3)** would be one valid value.
- **alpha**: controls the width of the network. This is known as the depth multiplier in the MobileNetV3 paper, but the name is kept for consistency with MobileNetV1 in Keras.
 - If **alpha < 1.0**, proportionally decreases the number of filters in each layer.
 - If **alpha > 1.0**, proportionally increases the number of filters in each layer.
 - If **alpha == 1**, default number of filters from the paper are used at each layer.
- **minimalistic**: In addition to large and small models this module also contains so-called minimalistic models, these models have the same per-layer dimensions characteristic as MobilenetV3 however, they don't utilize any of the advanced blocks (squeeze-and-excite units, hard-swish, and 5x5 convolutions). While these models are less efficient on CPU, they are much more performant on GPU/DSP.
- **include_top**: Boolean, whether to include the fully-connected layer at the top of the network. Defaults to **True**.
- **weights**: String, one of **None** (random initialization), **"imagenet"** (pre-training on ImageNet), or the path to the weights file to be loaded.
- **input_tensor**: Optional Keras tensor (i.e. output of **layers.Input()**) to use as image input for the model.
- **pooling**: String, optional pooling mode for feature extraction when **include_top** is **False**.
 - **None** means that the output of the model will be the 4D tensor output of the last convolutional block.
 - **avg** means that global average pooling will be applied to the output of the last convolutional block, and thus the output of the model will be a 2D tensor.
 - **max** means that global max pooling will be applied.
- **classes**: Integer, optional number of classes to classify images into, only to be specified if **include_top** is **True**, and if no **weights** argument is specified.
- **dropout_rate**: fraction of the input units to drop on the last layer.
- **classifier_activation**: A **str** or callable. The activation function to use on the "top" layer. Ignored unless **include_top=True**. Set **classifier_activation=None** to return the logits of the "top" layer. When loading pretrained weights, **classifier_activation** can only be **None** or **"softmax"**.
- **include_preprocessing**: Boolean, whether to include the preprocessing layer (**Rescaling**) at the bottom of the network. Defaults to **True**.
- **name**: String, the name of the model.

Call arguments

- **inputs**: A floating point **numpy.array** or backend-native tensor, 4D with 3 color channels, with values in the range **[0, 255]** if **include_preprocessing** is **True** and in the range **[-1, 1]** otherwise.

Returns

A model instance.

[\[source\]](#)

MobileNetV3Large function

```
keras.applications.MobileNetV3Large(  
    input_shape=None,  
    alpha=1.0,  
    minimalistic=False,  
    include_top=True,  
    weights="imagenet",  
    input_tensor=None,  
    classes=1000,  
    pooling=None,  
    dropout_rate=0.2,  
    classifier_activation="softmax",  
    include_preprocessing=True,  
    name="MobileNetV3Large",  
)
```

Instantiates the MobileNetV3Large architecture.

Reference

- [Searching for MobileNetV3](#) (ICCV 2019)

The following table describes the performance of MobileNets v3:

MACs stands for Multiply Adds

Classification Checkpoint	MACs(M)	Parameters(M)	Top1 Accuracy	Pixel1 CPU(ms)
mobilenet_v3_large_1.0_224	217	5.4	75.6	51.2
mobilenet_v3_large_0.75_224	155	4.0	73.3	39.8
mobilenet_v3_large_minimalistic_1.0_224	209	3.9	72.3	44.1
mobilenet_v3_small_1.0_224	66	2.9	68.1	15.8
mobilenet_v3_small_0.75_224	44	2.4	65.4	12.8
mobilenet_v3_small_minimalistic_1.0_224	65	2.0	61.9	12.2

For image classification use cases, see [this page for detailed examples](#).

For transfer learning use cases, make sure to read the [guide to transfer learning & fine-tuning](#).

Note: each Keras Application expects a specific kind of input preprocessing. For MobileNetV3, by default input preprocessing is included as a part of the model (as a **Rescaling** layer), and thus `keras.applications.mobilenet_v3.preprocess_input` is actually a pass-through function. In this use case, MobileNetV3 models expect their inputs to be float tensors of pixels with values in the **[0-255]** range. At the same time, preprocessing as a part of the model (i.e. **Rescaling** layer) can be disabled by setting `include_preprocessing` argument to **False**. With preprocessing disabled MobileNetV3 models expect their inputs to be float tensors of pixels with values in the **[-1, 1]** range.

Arguments

- **input_shape**: Optional shape tuple, to be specified if you would like to use a model with an input image resolution that is not **(224, 224, 3)**. It should have exactly 3 inputs channels. You can also omit this option if you would like to infer input_shape from an input_tensor. If you choose to include both input_tensor and input_shape then input_shape will be used if they match, if the shapes do not match then we will throw an error. E.g. **(160, 160, 3)** would be one valid value.
- **alpha**: controls the width of the network. This is known as the depth multiplier in the MobileNetV3 paper, but the name is kept for consistency with MobileNetV1 in Keras.
 - If **alpha < 1.0**, proportionally decreases the number of filters in each layer.
 - If **alpha > 1.0**, proportionally increases the number of filters in each layer.
 - If **alpha == 1**, default number of filters from the paper are used at each layer.
- **minimalistic**: In addition to large and small models this module also contains so-called minimalistic models, these models have the same per-layer dimensions characteristic as MobilenetV3 however, they don't utilize any of the advanced blocks (squeeze-and-excite units, hard-swish, and 5x5 convolutions). While these models are less efficient on CPU, they are much more performant on GPU/DSP.

- **include_top**: Boolean, whether to include the fully-connected layer at the top of the network. Defaults to `True`.
- **weights**: String, one of `None` (random initialization), `"imagenet"` (pre-training on ImageNet), or the path to the weights file to be loaded.
- **input_tensor**: Optional Keras tensor (i.e. output of `layers.Input()`) to use as image input for the model.
- **pooling**: String, optional pooling mode for feature extraction when `include_top` is `False`.
 - `None` means that the output of the model will be the 4D tensor output of the last convolutional block.
 - `avg` means that global average pooling will be applied to the output of the last convolutional block, and thus the output of the model will be a 2D tensor.
 - `max` means that global max pooling will be applied.
- **classes**: Integer, optional number of classes to classify images into, only to be specified if `include_top` is `True`, and if no `weights` argument is specified.
- **dropout_rate**: fraction of the input units to drop on the last layer.
- **classifier_activation**: A `str` or callable. The activation function to use on the "top" layer. Ignored unless `include_top=True`. Set `classifier_activation=None` to return the logits of the "top" layer. When loading pretrained weights, `classifier_activation` can only be `None` or `"softmax"`.
- **include_preprocessing**: Boolean, whether to include the preprocessing layer (`Rescaling`) at the bottom of the network. Defaults to `True`.
- **name**: String, the name of the model.

Call arguments

- **inputs**: A floating point `numpy.array` or backend-native tensor, 4D with 3 color channels, with values in the range `[0, 255]` if `include_preprocessing` is `True` and in the range `[-1, 1]` otherwise.

Returns

A model instance.
